

The Essential Contributions of Quantum Computing to Finance

A Pedagogical Journey Through Portfolio Optimization, Governance-First Shor Workflows,
Quantum Neural Networks, and Agentic Financial Discovery

Alejandro Reynoso

June 2026

Educational Manuscript Series in Quantum and Quantum-Inspired Finance

Companion Synthesis Paper for Four Research Manuscripts Developed at the Judge Business
School, University of Cambridge

Abstract

Quantum computing is often introduced to financial audiences through claims about future speedups, cryptographic disruption, or hardware breakthroughs. This manuscript adopts a different perspective. Its central argument is that the most durable contribution of quantum computing to finance may be representational rather than purely computational. By synthesizing four companion studies developed at the Judge Business School, University of Cambridge, together with one conceptual primer on quantum state-space reasoning, this paper presents quantum computational thinking as a structured learning pathway for financial professionals. The progression begins with a simple root-finding concept used to introduce quantum states and superposition, moves through portfolio optimization and Hamiltonian representations, examines governance-first Shor workflows as examples of problem transformation, introduces hybrid quantum neural networks for representation learning, and culminates in an agentic DQN-DNN-QNN architecture for financial discovery. Across these studies, the recurring lesson is that difficult financial problems often become easier to analyze when they are reformulated in representations that expose hidden structure. The manuscript is educational rather than experimental: it does not claim quantum advantage, propose new algorithms, or assert near-term production readiness. Instead, it provides a coherent conceptual framework for understanding how quantum and quantum-inspired methods can reshape thinking about optimization, learning, forecasting, governance, and discovery in finance.

Keywords: Quantum Computing; Quantum Finance; Portfolio Optimization; Shor’s Algorithm; Quantum Neural Networks; Quantum Machine Learning; Agentic AI; Reinforcement Learning; Computational Finance; Financial Technology.

Contents

1	Introduction	5
2	Research Program Overview	8
3	Conceptual Primer: A Simple Quantum Approach to Finding the Root of a Function	11
3.1	Motivation	12
3.2	Encoding Candidate Solutions into Quantum States	12
3.3	Superposition and Search	13
3.4	Interference and Measurement	13
3.5	Educational Lessons Learned	14
4	Quantum Portfolio Optimization	15
4.1	The Classical Portfolio Optimization Problem	15
4.2	Mapping Portfolio Selection into a Hamiltonian	16
4.3	Energy Minimization and Quantum Search	16
4.4	QAOA Intuition	17
4.5	Financial Interpretation of Results	17
5	A Governance-First Agentic Method for Teaching and Implementing Shor’s Algorithm: From Modular Period Finding to Auditable Quantum Workflow	18
5.1	Why Shor’s Algorithm Required an Agentic Explanation	18
5.2	From Factorization to Order Finding	19
5.3	The Agentic Decomposition of the Algorithm	19
5.4	Control Register and Modular Function Agents	20
5.5	Joint Register Construction and Entanglement	21
5.6	Target Measurement and Periodic Collapse	21
5.7	Fourier Analysis and Frequency Extraction	21
5.8	Control Measurement and Period Recovery	22
5.9	Factor Extraction and Final Explanation	22
5.10	Governance-First Interpretation	23
5.11	Why This Matters for Finance	23
6	Hybrid Quantum Neural Networks for Financial Prediction	24
6.1	Learning as Representation	24
6.2	Classical Neural Networks and Financial Forecasting	25
6.3	Parameterized Quantum Circuits	25
6.4	Quantum Feature Representations	26
6.5	Training and Hybrid Optimization	26

6.6	Experimental Insights and Educational Value	27
7	Agentic DQN-DNN-QNN Architectures for Financial Discovery	27
7.1	From Prediction to Discovery	28
7.2	Reinforcement Learning Foundations	28
7.3	Deep Q-Networks and Financial State Spaces	29
7.4	Quantum Neural Networks as Representation Engines	29
7.5	The Agentic Layer	30
7.6	Governance and Scientific Discovery	30
7.7	A Unified Framework for Financial Intelligence	31
8	Synthesis: A Unified Perspective on Quantum Computational Thinking	32
9	Companion Computational Laboratories	36
9.1	Conceptual Primer: A Simple Quantum Approach to Finding the Root of a Function	36
9.2	Companion Study 1: Quantum Portfolio Optimization	36
9.3	Companion Study 2: A Governance-First Agentic Method for Teaching and Implementing Shor's Algorithm: From Modular Period Finding to Auditable Quantum Workflow	36
9.4	Companion Study 3: Hybrid Quantum Neural Networks for Financial Prediction . .	37
9.5	Companion Study 4: An Agentic DQN-DNN-QNN Architecture for Financial Discovery	37
9.6	Additional Computational Laboratory: Agentic Quantum Optimization	38
9.7	Recommended Study Sequence	38
10	Limitations and Future Directions	38
10.1	An Educational Rather Than Experimental Contribution	39
10.2	Simplified Computational Environments	39
10.3	The Current State of Quantum Hardware	40
10.4	No Claim of Quantum Advantage	40
10.5	Interpretability and Governance	41
10.6	Future Directions	41
10.7	A Final Reflection on Scope	42
11	Conclusion	42
A	Concept Map of Quantum Computational Thinking	45
A.1	The Central Question	46
A.2	Level 1: Representation	46
A.3	Level 2: Optimization	47
A.4	Level 3: Transformation	47

A.5	Level 4: Learning	47
A.6	Level 5: Discovery	48
A.7	The Unified Framework	49
A.8	The Relationship Between the Conceptual Primer and Four Companion Studies . . .	49
A.9	The Essential Contributions of Quantum Computing to Finance	50
B	Glossary of Quantum Computing Terms for Financial Professionals	51
C	Educational Roadmap for Readers	54
C.1	The Educational Philosophy	54
C.2	Stage One: Learning to Think in State Spaces	55
C.3	Stage Two: Learning to Think in Optimization Landscapes	55
C.4	Stage Three: Learning the Power of Transformation	56
C.5	Stage Four: Learning Through Representation	57
C.6	Stage Five: Learning Through Discovery	57
C.7	The Complete Journey	58
C.8	Recommended Reading Sequence	58
C.9	Recommended Notebook Sequence	59
C.10	The Final Learning Objective	59

1 Introduction

The rapid emergence of quantum computing has generated considerable excitement across science, engineering, technology, and finance. Governments are investing billions of dollars in quantum research programs. Technology companies are racing to develop increasingly capable quantum hardware. Financial institutions are exploring potential applications in portfolio optimization, risk management, machine learning, forecasting, cryptography, and algorithmic trading. At the same time, an expanding academic literature continues to investigate the theoretical and practical implications of quantum computational methods.

Despite this growing interest, quantum computing remains a difficult subject for many financial professionals. Most introductions to the field assume significant familiarity with physics, mathematics, computer science, or information theory. Concepts such as Hilbert spaces, unitary transformations, eigenvectors, phase estimation, and quantum Fourier transforms often appear before readers have had an opportunity to develop an intuitive understanding of the underlying ideas. As a consequence, many practitioners encounter quantum computing as an abstract and highly technical discipline whose relevance to finance is difficult to assess.

This manuscript was written to address that gap.

The purpose of the present work is not to introduce new quantum algorithms, establish claims of quantum advantage, or present original scientific discoveries. Instead, its objective is educational. The manuscript synthesizes four companion studies developed at the Judge Business School, University of Cambridge, during June 2026, and uses a simple root-finding concept as an unsupported conceptual primer. Together, the primer and the four supported studies form a structured learning pathway designed to help financial professionals understand the most important conceptual contributions of quantum computing and quantum-inspired computation.

The central premise of this collection is straightforward.

The most valuable contribution of quantum computing to finance may not ultimately be faster computation. It may instead be the introduction of new representational frameworks through which complex financial problems can be understood, transformed, and analyzed.

This perspective differs from many popular discussions of quantum technologies. Public conversations frequently focus on computational speed, cryptographic disruption, or future hardware capabilities. While these topics are important, they can obscure a deeper lesson. Throughout the history of science, major advances have often emerged not from performing existing calculations more quickly, but from discovering better ways of representing problems.

The history of finance provides numerous examples.

Portfolio theory transformed investing by representing risk through covariance structures. Factor models transformed asset pricing by representing securities through common sources of risk. Option pricing became tractable through stochastic representations of uncertainty. Machine learning derives much of its power from its ability to construct useful representations of data.

Quantum computing belongs to this broader tradition.

Rather than representing information exclusively through deterministic binary states, quantum

computation introduces a richer representational framework based on amplitudes, phases, superposition, and interference. Whether or not future quantum computers eventually outperform classical systems on a wide range of practical financial tasks, these concepts already provide valuable intellectual tools for thinking about optimization, search, learning, forecasting, and decision-making.

For financial professionals, this perspective is particularly relevant because finance itself is fundamentally concerned with navigating large spaces of possibilities.

Portfolio managers search among thousands of potential allocations.

Risk managers evaluate enormous numbers of possible future scenarios.

Researchers search for patterns hidden within noisy and evolving datasets.

Traders continuously evaluate competing hypotheses regarding future market behavior.

In every case, the challenge involves identifying structure within complexity.

Quantum computational thinking offers a distinctive approach to this challenge.

The objective is not necessarily to replace classical methods. Rather, it is to explore alternative ways of representing problems so that useful structures become easier to identify. In many cases, the transformation of a problem may be as important as the eventual computation itself.

This idea serves as the organizing principle of one conceptual primer and four companion studies summarized throughout this manuscript.

The manuscript begins with a simple quantum approach to finding the root of a function. This material is included as a conceptual primer only; it does not correspond to a separate supporting paper in the companion collection. Although intentionally elementary, the example establishes the foundational concepts upon which the remainder of the collection is built. Readers are introduced to quantum states, superposition, amplitudes, interference, and measurement. More importantly, they begin to see computation as the evolution of a state within a structured space of possibilities.

The first supported companion study applies these ideas to portfolio optimization. Portfolio construction is transformed into a Hamiltonian representation in which investment decisions become candidate states within an energy landscape. The study demonstrates how quantum-inspired optimization frameworks encourage practitioners to think differently about allocation problems while preserving the economic intuition underlying classical portfolio theory.

The second supported companion study examines Shor’s algorithm, one of the most celebrated achievements in quantum computing. Beyond its implications for cryptography, Shor’s algorithm provides a remarkable illustration of representational transformation. Rather than solving factorization directly, the algorithm transforms the problem into one of period detection and then extracts the relevant information through phase estimation and the Quantum Fourier Transform. The study therefore highlights one of the most profound lessons of quantum computational thinking: difficult problems often become manageable when expressed within more revealing representations.

The third supported companion study introduces hybrid quantum neural networks. Modern machine learning systems derive much of their effectiveness from representation learning, and quantum neural networks extend this idea into quantum state spaces. The study explores how parameterized quantum circuits can be combined with classical deep learning architectures to create hybrid

models capable of constructing alternative feature representations for financial prediction tasks.

The fourth and final supported companion study culminates in an agentic discovery architecture integrating reinforcement learning, deep neural networks, quantum neural networks, and governance-first design principles. The objective is not merely prediction but discovery itself. Optimization, learning, representation, evaluation, and governance are combined into a structured framework capable of supporting autonomous financial research processes.

Viewed independently, these studies address very different computational problems.

Viewed collectively, they reveal a coherent intellectual progression.

The collection begins with quantum states.


It advances to optimization.

Optimization leads naturally to transformation.

Transformation supports learning.

Learning ultimately becomes a component of discovery.

The progression can be summarized as follows:



This progression represents the conceptual architecture of the entire manuscript.

The common thread connecting the conceptual primer and all four supported studies is the role of representation. Root finding becomes state-space exploration. Portfolio construction becomes energy minimization. Factorization becomes period detection. Prediction becomes representation learning. Research itself becomes an agentic process of discovery. In each case, the problem is reformulated into a representation that exposes useful structure.

For this reason, the present manuscript should not be viewed as a textbook on quantum mechanics, nor as a technical treatise on quantum algorithms. It is instead a pedagogical bridge between quantum computing and financial practice. Its purpose is to help readers develop intuition regarding the most important ideas that quantum computational thinking contributes to modern finance.

The companion Colab notebooks play a central role in this educational strategy. Each study is accompanied by an executable computational notebook that allows readers to experiment directly with the concepts under discussion. Rather than learning exclusively through mathematical derivation, readers are encouraged to observe, modify, and explore the behavior of quantum-inspired com-

putational models. Together, the papers and notebooks form an integrated learning environment designed to make quantum computing accessible to a broader audience of financial practitioners.

Ultimately, the objective of this collection is simple. By the conclusion of the manuscript, readers should possess a clear understanding of how quantum computational ideas can be used to rethink optimization, learning, forecasting, and decision-making. More importantly, they should appreciate that the most enduring contribution of quantum computing may not be computational acceleration alone, but a new framework for representing and understanding complexity itself.

2 Research Program Overview

The four companion studies synthesized in this manuscript were not conceived as independent explorations of unrelated quantum algorithms. Rather, they were designed as components of a single educational program whose objective is to help financial professionals develop intuition regarding the most important conceptual contributions of quantum computing and quantum-inspired computation.

The motivation for organizing the material in this manner arose from a practical observation. Readers entering the field of quantum computing frequently encounter a fragmented landscape. One article discusses qubits and superposition. Another focuses on portfolio optimization. A third examines Shor’s algorithm. Yet another explores quantum machine learning. While each topic may be interesting in isolation, the broader intellectual connections between them often remain unclear.

As a consequence, many practitioners develop a collection of disconnected concepts rather than a coherent understanding of quantum computational thinking.

The present collection attempts to address this challenge by presenting the material as a structured progression. Each study introduces a specific conceptual contribution of quantum computation and prepares the foundation for the study that follows. The objective is not merely to teach individual algorithms, but to demonstrate how seemingly diverse topics are connected by a common computational philosophy.

That philosophy can be summarized in a single idea:

The most important contribution of quantum computing is the ability to transform complex problems into alternative representations where useful structure becomes easier to identify.

Each manuscript explores a different manifestation of this principle.

Before the four supported companion studies begin, the manuscript introduces *A Simple Quantum Approach to Finding the Root of a Function* as a conceptual primer. This primer has no separate supporting paper in the collection. Its purpose is to introduce the concept of the quantum state. The problem itself is intentionally simple. The educational objective is not numerical sophistication but conceptual clarity. Readers encounter the foundational ideas of superposition, amplitudes, interference, and measurement. More importantly, they are introduced to a new way of

thinking about computation. Rather than viewing computation solely as a sequence of arithmetic operations, the primer presents it as the controlled evolution of a state within a structured space of possibilities.

This shift in perspective is essential because it establishes the conceptual foundation for everything that follows.

Once readers understand the idea of a state space, they are prepared to explore optimization. The first supported companion study, *Quantum Portfolio Optimization: A Canonical Educational Framework*, demonstrates how a familiar financial problem can be reformulated within a Hamiltonian representation. Portfolio allocations become candidate states. Constraints become penalties within an energy landscape. Optimization becomes a search for low-energy configurations.

This transformation introduces the second major lesson of the collection: the representation of a problem often determines how easily it can be solved.

The study therefore serves as the bridge between state-space thinking and optimization thinking. Readers begin to appreciate that the power of quantum computational methods frequently arises not from changing the objective, but from changing the mathematical framework through which the objective is expressed.

The second supported companion study, *A Governance-First Agentic Method for Teaching and Implementing Shor’s Algorithm: From Modular Period Finding to Auditable Quantum Workflow*, extends this idea dramatically. Shor’s algorithm is often presented as a breakthrough in cryptography because of its implications for integer factorization. While this interpretation is correct, the educational value of the algorithm is even broader.

Shor’s algorithm demonstrates one of the most powerful ideas in computational science: transformation.

The algorithm does not solve factorization directly. Instead, it converts factorization into a problem of period detection. Period detection is subsequently transformed into phase estimation, and phase estimation becomes accessible through the Quantum Fourier Transform. The original problem gradually disappears and is replaced by a sequence of alternative representations that reveal previously hidden structure.

For many readers, this study marks the point at which quantum computation ceases to appear mysterious and instead begins to resemble a general methodology for problem reformulation.

The third supported companion study, *Hybrid Quantum Neural Networks for Financial Prediction*, shifts the discussion from optimization and transformation toward learning. Modern machine learning systems derive much of their effectiveness from representation learning. Neural networks succeed because they transform raw observations into increasingly informative feature spaces.

Quantum neural networks extend this principle by introducing quantum state spaces as environments for feature construction and representation learning. The study explores how classical deep learning architectures and parameterized quantum circuits can be combined within a unified framework capable of generating alternative representations of financial data.

This progression is deliberate. The collection moves from representing candidate solutions to


representing patterns within data.

The fourth and final supported companion study, *An Agentic DQN-DNN-QNN Architecture for Financial Discovery*, completes the educational journey. The previous studies introduced states, optimization, transformation, and learning as individual computational capabilities. The final study integrates these capabilities into a governance-first discovery architecture.

Reinforcement learning, deep neural networks, quantum neural networks, and agentic evaluation mechanisms are combined within a structured framework designed to support autonomous financial research. The focus is no longer on solving a single problem but on creating a process through which hypotheses can be generated, tested, challenged, refined, and evaluated.

In this sense, the final study introduces discovery itself as a computational activity.

When viewed collectively, the conceptual primer and four supported studies form a coherent intellectual progression:



This progression represents the conceptual architecture of the entire research program.

The conceptual primer asks how possibilities can be represented.

The first supported study asks how desirable possibilities can be identified.

The second supported study asks how difficult problems can be transformed into more revealing forms.

The third supported study asks how useful representations can be learned from data.

The fourth supported study asks how optimization, transformation, and learning can be integrated into autonomous discovery systems.

The progression is cumulative. Each study introduces concepts that become necessary for understanding the next. Readers who proceed through the collection sequentially therefore develop intuition gradually rather than confronting the full complexity of quantum computing all at once.

From this perspective, the conceptual primer and four supported studies should not be viewed primarily as examples of individual quantum algorithms. They should instead be understood as chapters within a broader educational narrative concerning the role of representation in computation.

The common thread connecting the entire collection is remarkably consistent.

Root finding becomes state-space exploration.

Portfolio construction becomes energy minimization.

Factorization becomes period detection.

Prediction becomes representation learning.

Research becomes agentic discovery.

In every case, the problem is transformed into a representation in which useful structure becomes more visible.

This observation leads directly to the central argument of the present manuscript. The enduring significance of quantum computing for financial professionals may not lie exclusively in future hardware advances or computational acceleration. Its deeper contribution may be the introduction of a new language for thinking about complexity.

The chapters that follow examine each stage of this progression in detail. Together with the companion Colab notebooks, they provide a practical pathway through which readers can move from elementary concepts such as superposition and measurement to advanced topics such as hybrid machine learning systems and agentic financial discovery architectures.

The objective is not to produce quantum physicists. The objective is to equip financial professionals with a conceptual framework through which they can understand how quantum computational thinking may influence the future of optimization, learning, forecasting, and decision-making.

3 Conceptual Primer: A Simple Quantum Approach to Finding the Root of a Function

This conceptual primer was intentionally designed to be simple. Before discussing portfolio optimization, cryptographic algorithms, quantum neural networks, or reinforcement learning architectures, it is necessary to develop an intuition for how quantum computation represents and manipulates information. For many newcomers, quantum computing appears intimidating because it introduces unfamiliar concepts such as superposition, interference, phase, and measurement. Yet these concepts become considerably easier to understand when they are attached to a problem that is already familiar.

The problem selected for this introductory study was the search for the root of a mathematical function. In classical mathematics, finding the root of a function means identifying the value of a variable that causes the function to evaluate to zero. This type of problem appears throughout science, engineering, economics, and finance. Equilibrium models, optimization procedures, calibration exercises, and forecasting systems frequently require the solution of equations whose roots must be determined numerically.

The objective of the primer was not to demonstrate a practical quantum advantage. Instead, it was to illustrate how a classical search problem can be reformulated within a quantum computational framework.

3.1 Motivation

Suppose that we wish to identify the value of a variable (x) such that

$$f(x) = 0.$$

A classical algorithm typically evaluates candidate values one at a time. Depending on the complexity of the function, it may use methods such as bisection, Newton-Raphson iterations, gradient-based optimization, or exhaustive search. Although these methods differ substantially, they all share a common characteristic: candidate solutions are evaluated sequentially.

Quantum computation encourages us to think differently.

Instead of representing a single candidate solution, a quantum state can represent a collection of candidate solutions simultaneously. The purpose of this example is therefore not to accelerate root finding but to introduce the idea that multiple possibilities may coexist within a single computational object.

This conceptual shift is perhaps the most important lesson of the entire study.

The quantum state becomes a container of possibilities.

Rather than examining one candidate at a time, the computational framework allows us to reason about many candidates simultaneously and manipulate their relationships through quantum operations.

3.2 Encoding Candidate Solutions into Quantum States

The first step in any quantum algorithm is representation.

A quantum computer cannot operate directly on abstract mathematical concepts. The problem must first be encoded into quantum states. In the root-finding example, candidate values of the variable are associated with basis states of a quantum register.

For example, a three-qubit system contains eight computational basis states:

$$|000\rangle, |001\rangle, |010\rangle, |011\rangle, |100\rangle, |101\rangle, |110\rangle, |111\rangle.$$

These states may be interpreted as discrete candidate values for the variable under consideration.

The remarkable feature of quantum mechanics is that the system does not need to occupy only one of these states. Instead, it may exist in a superposition:

$$|\psi\rangle = a_0|000\rangle + a_1|001\rangle + \cdots + a_7|111\rangle.$$

The coefficients a_i are known as amplitudes.

Their squared magnitudes determine the probabilities observed when the system is measured.

At this stage, the quantum state represents an entire collection of candidate solutions. The system has not yet selected a root. Rather, it contains information about all possible candidates simultaneously.

This idea is often described as quantum parallelism. While the term is sometimes overstated in popular discussions, it captures an important intuition. Quantum states allow information about multiple possibilities to coexist within a single mathematical object.

3.3 Superposition and Search

Once candidate solutions have been encoded into the quantum state, the next question becomes how to distinguish desirable candidates from undesirable ones.

In a classical search process, this distinction is achieved through direct evaluation. A candidate is tested, the function is computed, and a decision is made regarding whether to continue searching.

The quantum framework introduces a different perspective.

Instead of evaluating candidates independently, the algorithm modifies the amplitudes associated with the candidate states.

Initially, the amplitudes may be distributed uniformly across all possibilities. Every candidate receives equal weight. As transformations are applied, some amplitudes increase while others decrease.

The resulting state evolves according to unitary operators.

Geometrically, one may think of the state vector as rotating within a high-dimensional space. The purpose of these rotations is to concentrate probability mass around states corresponding to desirable solutions.

Although our introductory example remains deliberately simple, it introduces an idea that will reappear throughout this collection: computation as state evolution.

In classical computing, algorithms are frequently viewed as sequences of instructions.

In quantum computing, algorithms are often better understood as sequences of transformations applied to vectors.

The distinction may seem subtle, but it represents a fundamentally different computational philosophy.

3.4 Interference and Measurement

Perhaps the most important concept introduced by this study is interference.

When financial professionals first encounter quantum computing, they often focus exclusively on probabilities. This is understandable because probability theory already plays a central role in finance. However, probabilities alone cannot explain quantum computation.

The key ingredient is phase.

Each amplitude possesses both magnitude and phase. Two states may have identical probabilities while exhibiting different phase relationships.

These phase relationships determine how amplitudes combine.

When amplitudes reinforce one another, constructive interference occurs. When amplitudes cancel one another, destructive interference occurs.

This mechanism provides the foundation for many quantum algorithms.

Instead of directly selecting a solution, the algorithm manipulates phase relationships so that desirable states receive constructive reinforcement while undesirable states experience cancellation.

The probability distribution observed at the end of the computation is therefore not arbitrary. It is the result of carefully engineered interference patterns.

For many readers, this is the moment at which quantum computation begins to feel genuinely different from classical computation.

The algorithm is no longer simply evaluating alternatives.

It is shaping a probability landscape through the manipulation of phase information.

Measurement then converts this quantum state into a classical outcome.

When the system is observed, the superposition collapses into one of the computational basis states. The probability of observing a particular state depends on the amplitude structure created during the computational process.

In practical terms, measurement extracts information from the transformed state.

3.5 Educational Lessons Learned

Although the root-finding example is mathematically simple, it serves a critical educational purpose within the overall collection.

First, it introduces the notion that computation begins with representation. Before a problem can be solved, it must be encoded into a state space. This principle remains central throughout every subsequent study.

Second, it establishes the importance of superposition. Rather than treating candidate solutions independently, quantum representations allow collections of possibilities to be considered simultaneously.

Third, it introduces interference as a computational mechanism. Unlike classical algorithms, which typically rely on explicit evaluation and comparison, quantum algorithms often derive their power from the constructive and destructive combination of amplitudes.

Fourth, it encourages a geometric interpretation of computation. State vectors evolve through transformations that may be viewed as rotations within abstract spaces. This geometric perspective becomes increasingly important when discussing optimization algorithms, Fourier transforms, and quantum neural networks.

Finally, the study demonstrates that quantum computation is fundamentally about changing perspective. The objective is not merely to solve a problem but to reformulate the problem within a representation that reveals useful structure.

This lesson provides the conceptual bridge to the next study in the collection. Once readers understand how a simple search problem can be represented within a quantum state space, it becomes possible to examine more realistic applications. Portfolio optimization offers a natural next step because it transforms a familiar financial decision problem into a search over a much larger collection of candidate states. The ideas introduced here—representation, superposition,

interference, and measurement—will therefore reappear repeatedly as we move from elementary examples toward increasingly sophisticated financial applications.

4 Quantum Portfolio Optimization

Portfolio optimization represents one of the most natural meeting points between finance and quantum computation. Unlike cryptographic applications, which may appear distant from the daily activities of most investment professionals, portfolio construction is a problem that lies at the very center of modern financial practice. Asset managers, pension funds, insurance companies, sovereign wealth funds, endowments, and individual investors all face the same fundamental challenge: how should capital be allocated across a collection of available investment opportunities?

This challenge becomes increasingly difficult as the number of assets grows. Even relatively simple allocation problems can generate extraordinarily large search spaces. The introduction of constraints such as diversification requirements, sector limits, turnover restrictions, liquidity thresholds, regulatory conditions, and risk budgets further increases the complexity of the problem. For this reason, portfolio optimization provides an ideal educational setting in which to explore quantum computational thinking.

The objective of this study was not merely to apply a quantum algorithm to a financial problem. Rather, it was to demonstrate how a familiar investment decision can be transformed into a representation that naturally aligns with quantum optimization techniques.

4.1 The Classical Portfolio Optimization Problem

The foundations of modern portfolio theory can be traced to the pioneering work of Harry Markowitz. The central idea is deceptively simple. Investors seek portfolios that maximize expected return while simultaneously minimizing risk. Risk is commonly represented through the covariance structure of asset returns, reflecting the reality that securities do not move independently.

In the classical formulation, portfolio optimization is often expressed as a constrained mathematical program:

$$\max_w \left(\mu^T w - \lambda w^T \Sigma w \right),$$

where w represents portfolio weights, μ denotes expected returns, Σ is the covariance matrix, and λ controls the trade-off between return and risk.

Although elegant, this formulation can become computationally challenging when realistic constraints are added. The number of possible portfolio combinations grows rapidly as the number of securities increases. Furthermore, many practical investment decisions involve discrete choices rather than continuous allocations. An investor may need to decide whether to include or exclude particular securities, select among competing strategies, or satisfy complex allocation rules.

From a computational perspective, portfolio optimization becomes a search problem.

The challenge is to identify the portfolio configuration that best satisfies the investor’s objectives while respecting all relevant constraints.

This observation creates a natural bridge to quantum optimization.

4.2 Mapping Portfolio Selection into a Hamiltonian

The most important conceptual step in the study is the transformation of the portfolio optimization problem into a Hamiltonian representation.

In physics, a Hamiltonian describes the energy of a system. Systems naturally evolve toward low-energy configurations. Quantum optimization methods exploit this principle by encoding a computational objective into an energy landscape.

From a financial perspective, the interpretation is straightforward.

Rather than searching directly for an optimal portfolio, we construct an energy function whose minimum corresponds to the desired allocation.

Expected returns, risk penalties, diversification constraints, and investment restrictions can all be incorporated into the Hamiltonian. Portfolios that violate constraints receive energy penalties, while portfolios that satisfy objectives receive lower energy values.

The optimization problem is therefore transformed into a search for low-energy states.

This transformation represents a powerful example of the broader theme discussed throughout this paper: representation matters.

The underlying financial problem has not changed.

The investor still seeks an attractive portfolio.

However, the mathematical description of the problem has been completely reformulated. What was originally an investment decision problem has become an energy minimization problem.

This change of perspective allows quantum computational tools to be applied naturally.

4.3 Energy Minimization and Quantum Search

Once the Hamiltonian has been defined, candidate portfolios can be represented as basis states within a quantum register.

Consider a simple example involving a collection of assets. Each basis state may correspond to a particular portfolio configuration. Some states represent desirable allocations, while others correspond to portfolios that are excessively risky, poorly diversified, or inconsistent with investment constraints.

The quantum state is initialized as a superposition over many possible portfolios.

Instead of evaluating allocations one by one, the computational process begins with a representation that simultaneously encompasses a large collection of alternatives.

The objective is then to modify the quantum state so that favorable portfolios become increasingly prominent.

From a geometric perspective, the state vector evolves through a sequence of transformations. The amplitudes associated with attractive allocations increase, while the amplitudes associated with inferior allocations decrease.

This process introduces a fundamentally different way of thinking about optimization.

Classical optimization often emphasizes explicit comparisons among alternatives. Quantum optimization focuses instead on shaping a probability distribution over the space of possible solutions.

The final measurement reveals a portfolio whose probability of selection reflects the structure imposed by the optimization procedure.

Although practical implementations remain constrained by current hardware limitations, the educational value of this framework is substantial. It encourages financial professionals to view optimization not simply as a search over numerical values but as a transformation of a state space.

4.4 QAOA Intuition

The principal computational framework explored in this study was inspired by the Quantum Approximate Optimization Algorithm (QAOA).

QAOA occupies an important position within quantum computing because it combines conceptual simplicity with practical relevance. Rather than requiring large fault-tolerant quantum computers, it is designed to operate within hybrid quantum-classical environments.

At a high level, QAOA alternates between two operators.

The first operator reflects the optimization objective and is derived from the Hamiltonian. This operator introduces information about portfolio quality into the quantum state.

The second operator acts as a mixing mechanism. Its purpose is to ensure that the system continues to explore the solution space rather than becoming trapped within a limited region.

The repeated application of these operators gradually reshapes the state vector.

An intuitive analogy may be useful. Imagine a landscape containing hills and valleys. The Hamiltonian identifies where the valleys are located, while the mixing operator allows the system to move across the landscape. Through successive iterations, the state becomes increasingly concentrated around attractive regions.

From the perspective of a financial practitioner, the most important insight is that optimization emerges from the interaction between objective information and controlled exploration.

This idea is not entirely foreign to finance. Many portfolio construction methods balance exploitation of known opportunities against exploration of alternative allocations. QAOA formalizes a similar balance within a quantum computational framework.

4.5 Financial Interpretation of Results

The most important contribution of this study lies not in demonstrating superior investment performance but in illustrating how financial problems can be reimaged through quantum representations.

Several educational lessons emerge.

First, portfolio optimization can be interpreted as a search over a state space rather than a collection of numerical calculations. This perspective encourages practitioners to think geometrically about investment decisions.

Second, constraints become an integral part of the representation itself. Instead of being treated as external conditions imposed after optimization, they are incorporated directly into the energy landscape.

Third, the Hamiltonian framework reveals a deep connection between physics and finance. Both disciplines seek configurations that satisfy competing objectives under complex constraints. While the interpretations differ, the mathematical structures often exhibit remarkable similarities.

Finally, the study reinforces the central theme of this collection. Quantum computational thinking is fundamentally about transformation. The investor’s problem remains unchanged. Capital must still be allocated, risk must still be controlled, and returns must still be pursued. What changes is the representation through which the problem is examined.

This insight provides a natural transition to the next study. Whereas portfolio optimization transforms allocation into energy minimization, Shor’s algorithm demonstrates an even more dramatic transformation. In that case, the challenge of integer factorization is reformulated as a problem of period detection. The resulting algorithm provides one of the most striking examples of how changing representations can reveal hidden structure and create entirely new computational possibilities.

5 A Governance-First Agentic Method for Teaching and Implementing Shor’s Algorithm: From Modular Period Finding to Auditable Quantum Workflow

5.1 Why Shor’s Algorithm Required an Agentic Explanation

Shor’s algorithm presented one of the most difficult pedagogical challenges in the collection. Traditional explanations are mathematically correct, but they often compress too many ideas into a single circuit narrative: modular arithmetic, two-register construction, reversible exponentiation, entanglement, measurement, periodic collapse, the Quantum Fourier Transform, frequency sampling, rational reconstruction, and classical factor extraction.

For financial professionals and non-quantum specialists, this compression can make the algorithm appear opaque. The learner may understand that Shor’s algorithm factors integers by finding a period, but still fail to understand what is periodic, where the periodicity is encoded, why two registers are required, why measurement appears in the middle of the algorithm, why the Quantum Fourier Transform is necessary, and why classical arithmetic remains essential at the end.

For that reason, the companion study did not merely present Shor’s algorithm as a circuit. It introduced a governance-first agentic method for teaching and implementing the algorithm. The agentic architecture decomposes the workflow into specialized computational agents, each with

explicit responsibilities, inputs, outputs, validation rules, and audit evidence.

This was the pedagogical breakthrough.

The algorithm became understandable because it was no longer presented as a monolithic quantum circuit. It became a sequence of auditable transformations.

5.2 From Factorization to Order Finding

The central mathematical idea remains the standard reduction used in Shor's algorithm. The algorithm does not search directly for the factors of an integer N . Instead, it chooses a base a such that $\gcd(a, N) = 1$, and studies the modular function

$$f(x) = a^x \bmod N.$$

The objective is to find the smallest positive integer r such that

$$a^r \equiv 1 \pmod{N}.$$

This integer r is the order of a modulo N . Once a valid even period is found, and provided that $a^{r/2} \not\equiv -1 \pmod{N}$, the factors can be recovered using

$$\gcd(a^{r/2} - 1, N) \quad \text{and} \quad \gcd(a^{r/2} + 1, N).$$

For the canonical educational example $N = 15$ and $a = 2$, the modular sequence is

$$1, 2, 4, 8, 1, 2, 4, 8, \dots$$


so the period is ($r=4$). The factors are then recovered as


$$\gcd(2^2 - 1, 15) = 3, \quad \gcd(2^2 + 1, 15) = 5.$$

The important pedagogical point is that Shor's algorithm is not primarily about trying possible factors. It is about transforming factorization into order finding.

5.3 The Agentic Decomposition of the Algorithm

The companion notebook decomposes the algorithm into a governed workflow of specialized agents:





Each agent performs one bounded conceptual task. This separation is essential. The Control Register Agent is not allowed to alter the modular function. The Modular Function Agent is not allowed to infer the period from Fourier samples. The Period Recovery Agent is not allowed to read the benchmark period. The LLM Explanation Agent can explain the audit trail, but it cannot alter the computational record.

This architecture transforms the notebook into an auditable educational workflow. Every stage has declared inputs, declared outputs, validation checks, and recorded evidence.

5.4 Control Register and Modular Function Agents

The Control Register Agent prepares the control register. With n control qubits, the register has dimension

$$Q = 2^n.$$

Applying Hadamard gates creates the equal superposition

$$\frac{1}{\sqrt{Q}} \sum_{x=0}^{Q-1} |x\rangle.$$

This register represents the coherent input domain. It is not a classical list of integers. It is a quantum state whose basis labels correspond to possible exponents.

The Modular Function Agent evaluates the function

$$a^x \bmod N$$

over the relevant input range. In the educational notebook, the classical order may be computed as a benchmark, but this benchmark is not passed to the recovery agent. This prevents answer leakage and preserves the integrity of the workflow.

5.5 Joint Register Construction and Entanglement

The Joint Register Agent constructs the joint state

$$\frac{1}{\sqrt{Q}} \sum_{x=0}^{Q-1} |x\rangle |a^x \bmod N\rangle.$$

This equation is central to the paper. It clarifies a common misunderstanding: the modular values are not amplitudes. They are labels in the target register.

For $N = 15$ and $a = 2$, the first branches are

$$|0\rangle|1\rangle + |1\rangle|2\rangle + |2\rangle|4\rangle + |3\rangle|8\rangle + |4\rangle|1\rangle + \dots$$

The repeated modular values encode periodicity. Entanglement arises because the target value depends on the control value. It is not an arbitrary resource added to the circuit; it is the structural consequence of coherent modular exponentiation.

5.6 Target Measurement and Periodic Collapse

The Target Measurement Agent measures the second register. If the measured modular value is y , the joint state collapses onto the set of control values satisfying

$$a^x \bmod N = y.$$

For example, if $y = 1$, the surviving control values are

$$0, 4, 8, 12, \dots$$

and the control register becomes a periodic coset state:

$$\frac{1}{\sqrt{M}} (|0\rangle + |4\rangle + |8\rangle + |12\rangle + \dots).$$

This is one of the most important conceptual moments in the algorithm. Measurement does not reveal the period directly. It creates a state whose support is periodic.

The Periodic Pattern Agent inspects this structure for educational purposes. This inspection helps the learner see that the hidden order is now encoded in the spacing of the surviving control-register amplitudes.

5.7 Fourier Analysis and Frequency Extraction

At this stage, measuring the control register directly would return only one surviving position. That would not efficiently reveal the spacing.

The Fourier Agent applies the Quantum Fourier Transform:

$$QFT_Q|x\rangle = \frac{1}{\sqrt{Q}} \sum_{k=0}^{Q-1} e^{2\pi i x k/Q} |k\rangle.$$

The QFT converts regular spacing in the computational basis into concentration in the frequency basis. For a periodic state of the form

$$\frac{1}{\sqrt{M}} \sum_{m=0}^{M-1} |x_0 + mr\rangle,$$

the Fourier amplitudes concentrate near values satisfying

$$k \approx \frac{jQ}{r}.$$

The QFT does not output the period directly. It creates a probability distribution in which frequency-related outcomes are more likely.

This is the core quantum mechanism: phase relationships become observable probability differences through interference.

5.8 Control Measurement and Period Recovery

The Control Measurement Agent samples the transformed control register. A measured value (k) is interpreted through

$$\frac{k}{Q} \approx \frac{j}{r}.$$

The Period Recovery Agent then uses rational approximation or continued fractions to generate candidate periods. Importantly, each candidate must be validated by checking

$$a^r \equiv 1 \pmod{N}.$$

This validation step converts a probabilistic measurement clue into a verified arithmetic statement. Failed candidates are not hidden. They remain part of the audit trail.

5.9 Factor Extraction and Final Explanation

Once a valid period is recovered, the Factor Extraction Agent verifies that the period is usable. It must be even, and it must satisfy the nontriviality condition. The agent then computes

$$\gcd(a^{r/2} - 1, N) \quad \text{and} \quad \gcd(a^{r/2} + 1, N).$$

The resulting factors are accepted only if they are nontrivial and multiply back to N .

Finally, the LLM Explanation Agent receives the governed evidence: configuration, measurements, candidate periods, validation checks, factors, and audit log. Its role is explanatory, not

computational. It narrates what happened, distinguishes the NumPy simulation from scalable hardware implementation, and explains success or failure modes.

5.10 Governance-First Interpretation

The agentic architecture is not cosmetic. It changes the way Shor’s algorithm is taught, inspected, and validated.

Each stage becomes a computational claim:

- The control state was prepared correctly.
- The modular map was computed correctly.
- The joint register preserved probability.
- The target measurement produced a valid collapse.
- The surviving control support was periodic.
- The Fourier transform preserved normalization.
- The measured frequency suggested a candidate period.
- The candidate period passed modular validation.
- The extracted factors were nontrivial and correct.

This is why the paper is governance-first. It treats quantum computation not as an opaque circuit but as a chain of auditable claims.

5.11 Why This Matters for Finance

The broader significance of the study is not limited to factorization. The same agentic governance pattern applies naturally to quantum-financial workflows.

A quantum portfolio optimizer must explain how returns, covariance matrices, constraints, penalties, and bitstring outputs are encoded and validated.

An amplitude-estimation pricing routine must explain how scenarios, payoff functions, amplitudes, samples, and confidence intervals are constructed.

A quantum risk engine must document data lineage, encoding assumptions, sampling uncertainty, benchmark comparison, and failure conditions.

The Shor notebook therefore becomes both a pedagogical implementation and a prototype for governable hybrid quantum-classical systems.

The central lesson is methodological:

Quantum algorithms become more understandable and more governable when they are represented not as opaque monolithic circuits, but as sequences of explicit, auditable decisions.

That lesson is directly relevant to finance, where transparency, validation, auditability, and model governance are not optional additions but essential requirements.

6 Hybrid Quantum Neural Networks for Financial Prediction

The conceptual primer and previous supported studies focused primarily on search, optimization, and transformation. The root-finding primer introduced the idea of representing candidate solutions within a quantum state space. Portfolio optimization demonstrated how investment decisions can be reformulated as energy minimization problems. Shor’s algorithm illustrated how an apparently difficult arithmetic problem can become tractable after being transformed into a problem of period detection. Although these applications differ substantially, they share a common objective: extracting useful structure from complex systems.

Machine learning pursues a similar goal.

Modern predictive models seek to discover patterns hidden within large collections of data. Financial markets provide a particularly challenging environment because signals are often weak, noisy, nonlinear, and constantly evolving. Traditional statistical methods and deep learning architectures have achieved considerable success in forecasting, classification, anomaly detection, and portfolio construction. Nevertheless, financial prediction remains one of the most difficult problems in applied data science.

The purpose of the third supported study was therefore to investigate whether quantum computational ideas could contribute to the process of representation learning. More specifically, the study explored a hybrid architecture combining classical convolutional neural networks with parameterized quantum circuits. The objective was not to replace classical machine learning but rather to examine how quantum representations might enrich the information available to predictive systems.

6.1 Learning as Representation

One of the most important developments in modern artificial intelligence has been the recognition that successful prediction depends heavily on representation.

Historically, machine learning practitioners devoted enormous effort to feature engineering. Domain experts manually constructed variables that they believed would capture useful patterns in the data. In finance, such features might include moving averages, momentum indicators, volatility measures, valuation ratios, yield spreads, or macroeconomic variables.

Deep learning altered this paradigm.

Instead of requiring researchers to specify all relevant features explicitly, neural networks learned representations directly from data. Through successive layers of transformation, raw observations were converted into increasingly abstract and informative structures.

From a conceptual perspective, this process bears a striking resemblance to the transformations discussed in earlier sections of this paper.

Once again, the challenge is not merely computation.

The challenge is representation.

A neural network succeeds because it discovers representations that make prediction easier.

This observation creates a natural connection between machine learning and quantum computing. If quantum state spaces provide alternative representations of information, then it becomes reasonable to ask whether they can contribute to the learning process itself.

6.2 Classical Neural Networks and Financial Forecasting

The hybrid architecture examined in this study began with a conventional convolutional neural network (CNN).

Although CNNs are most commonly associated with image processing, their underlying principles extend naturally to many forms of structured data. Financial time series often exhibit local patterns, recurring structures, and relationships across neighboring observations. Convolutional filters can therefore be used to extract informative features from sequences of market data.

In the study, synthetic financial datasets were employed to maintain complete control over the experimental environment. Variables representing market conditions, asset characteristics, and predictive signals were processed through convolutional layers designed to identify relevant patterns.

The CNN served two important functions.

First, it reduced the dimensionality of the input data by extracting higher-level features.

Second, it created an interface between classical information and the quantum components of the model.

This hybrid structure reflects a practical reality of current quantum computing. Classical and quantum systems are likely to coexist for many years. Rather than viewing them as competing technologies, it is often more productive to treat them as complementary components within a larger computational architecture.

6.3 Parameterized Quantum Circuits

The distinctive element of the architecture was the inclusion of a parameterized quantum circuit, commonly referred to as a quantum neural network (QNN).

A QNN differs fundamentally from a classical neural network.

In a conventional network, information is transformed through matrix multiplications followed by nonlinear activation functions. The model learns by adjusting weights and biases to reduce prediction errors.

In a quantum neural network, information is encoded into quantum states. The state then evolves through a sequence of parameterized quantum gates whose rotation angles play a role analogous to trainable weights.

Each gate modifies the quantum state by changing amplitude distributions and phase relationships. The resulting state is subsequently measured, producing outputs that can be incorporated into the overall learning process.

This architecture introduces an important conceptual shift.

The network is no longer operating exclusively within a classical vector space.

Instead, part of the computation occurs within a quantum state space characterized by superposition, interference, and unitary evolution.

The model learns by adjusting parameters that govern the evolution of quantum states.

Although the underlying mathematics differs substantially from classical deep learning, the intuition remains familiar. The objective is to discover transformations that improve predictive performance.

6.4 Quantum Feature Representations

The central hypothesis of the study was that quantum state spaces may provide useful feature representations.

This hypothesis is often misunderstood.

The purpose is not necessarily to achieve a dramatic computational speedup. Rather, the objective is to investigate whether quantum transformations generate representations that are difficult to obtain through conventional architectures.

The key mechanism is once again phase.

As discussed in previous sections, quantum states contain information not only through probabilities but also through phase relationships. These phase relationships influence interference patterns, which in turn affect measurement outcomes.

From a machine learning perspective, one may view the quantum circuit as a highly structured feature transformation engine.

Input data are encoded into quantum states.

The circuit then applies a sequence of rotations and entangling operations.

These operations reshape the state within a high-dimensional Hilbert space.

Measurement converts the resulting state into classical outputs that can be processed by downstream layers.

The hope is that the quantum transformation uncovers useful structures that improve the model's ability to distinguish among different patterns in the data.

Whether this advantage exists universally remains an open research question. Nevertheless, the study demonstrates how quantum computational concepts can be integrated naturally into modern machine learning pipelines.

6.5 Training and Hybrid Optimization

A particularly interesting aspect of the hybrid architecture is the training process.

Unlike purely classical models, the optimization procedure must coordinate parameters across both classical and quantum components.

The CNN learns conventional weights through gradient-based optimization.

Simultaneously, the QNN learns rotation angles associated with quantum gates.

The training process therefore becomes a collaborative effort between classical and quantum representations.

This hybrid strategy reflects a broader theme within contemporary quantum computing. Rather than expecting quantum systems to solve entire problems independently, many practical approaches focus on combining classical strengths with quantum capabilities.

For financial applications, this perspective is especially attractive. Financial institutions already possess extensive investments in classical infrastructure, data pipelines, and machine learning systems. Hybrid architectures offer a pathway through which quantum methods can be explored without requiring a complete redesign of existing workflows.

6.6 Experimental Insights and Educational Value

The principal contribution of the study was educational rather than empirical.

The results demonstrated that quantum circuits can function as trainable computational layers embedded within larger predictive architectures. More importantly, they provided a concrete example of how quantum computational thinking extends beyond optimization and cryptography.

Several lessons emerged.

First, machine learning and quantum computing share a common emphasis on representation. Both disciplines seek transformations that reveal hidden structure within data.

Second, quantum state spaces offer a fundamentally different environment for constructing representations. Information is encoded not only through numerical values but also through amplitudes, phases, and interference effects.

Third, hybrid architectures provide a practical bridge between current machine learning practice and emerging quantum technologies.

Finally, the study reinforced a central theme of this collection: the value of quantum computation often lies in how problems are represented rather than merely how quickly they are solved.

This realization provides a natural transition to the final study in the collection. If quantum circuits can contribute to predictive modeling, then it becomes reasonable to ask whether they can also contribute to autonomous discovery systems. The next section explores this possibility through an agentic architecture that combines reinforcement learning, deep neural networks, and quantum neural networks within a governance-first framework for financial research and decision-making.

7 Agentic DQN-DNN-QNN Architectures for Financial Discovery

The conceptual primer and supported studies in this collection progressively expanded the role of quantum computational thinking. We began with a simple search concept, moved into portfolio

optimization, examined the transformation-based logic of Shor’s algorithm, and then explored how quantum state spaces could be incorporated into predictive machine learning systems. The natural next question is whether these ideas can be integrated into a broader framework capable of supporting autonomous financial research and decision-making.

This question motivated the fourth and final supported study.

While the earlier papers focused on specific computational tasks, this study considered the design of an entire discovery architecture. The objective was not merely to predict market outcomes or optimize a portfolio. Instead, the goal was to create a system capable of generating hypotheses, evaluating alternatives, learning from experience, and refining its behavior through repeated interaction with a financial environment.

To achieve this objective, the study combined three distinct computational paradigms: Deep Q-Networks (DQNs), Deep Neural Networks (DNNs), and Quantum Neural Networks (QNNs). These components were embedded within an agentic framework governed by explicit oversight mechanisms intended to ensure transparency, auditability, and scientific rigor.

The resulting architecture represented the culmination of the educational journey presented throughout this collection.

7.1 From Prediction to Discovery

Traditional machine learning systems are often designed to answer narrowly defined questions.

Will a stock rise or fall tomorrow?

Will a borrower default?

Will volatility increase?

These questions are important, but they represent only a small portion of the research process. Human analysts typically engage in a much broader sequence of activities. They formulate hypotheses, collect evidence, compare competing explanations, revise assumptions, and continuously adapt their understanding as new information becomes available.

In other words, financial research is fundamentally a discovery process.

The agentic architecture explored in this study was designed to mimic aspects of that process. Rather than viewing artificial intelligence as a static prediction engine, the framework treats it as a dynamic participant in an ongoing cycle of exploration and learning.

This distinction is important.

A predictive model seeks answers.

A discovery system seeks understanding.

While both approaches rely on data, the second operates at a higher level of abstraction. It attempts to generate new ideas rather than simply evaluate existing ones.

7.2 Reinforcement Learning Foundations

The primary learning mechanism employed in the architecture was reinforcement learning.

Unlike supervised learning, which relies on labeled examples, reinforcement learning is based on interaction with an environment. An agent observes a state, selects an action, receives feedback in the form of rewards, and updates its behavior accordingly.

Financial decision-making naturally fits this framework.

Portfolio managers allocate capital.

Traders enter and exit positions.

Risk managers impose constraints.

Each decision produces consequences that can be evaluated over time.

The reinforcement learning paradigm captures this sequential nature of decision-making.

At the heart of the architecture lies the concept of a policy: a mapping from observed states to actions. The objective of learning is to discover policies that maximize long-term rewards while managing uncertainty and risk.

This formulation introduces a powerful perspective. Rather than predicting isolated events, the system learns strategies.

The emphasis shifts from forecasting to decision-making.

7.3 Deep Q-Networks and Financial State Spaces

The first major component of the architecture was the Deep Q-Network.

A DQN approximates the value of actions using a neural network. Given a particular state of the environment, the network estimates the expected reward associated with available actions.

In financial applications, states may include market variables, technical indicators, macroeconomic conditions, portfolio characteristics, or alternative representations generated by upstream models.

The DQN transforms these observations into estimates of future value.

Conceptually, the network learns a map of the decision landscape.

Actions associated with favorable outcomes receive higher valuations, while less attractive actions receive lower valuations.

Over time, the system develops an increasingly refined understanding of which actions are likely to produce desirable results.

From an educational perspective, the DQN serves as the classical backbone of the architecture. It provides the framework through which decisions are evaluated and improved.

However, classical representations are not the only source of information available to the system.

7.4 Quantum Neural Networks as Representation Engines

The second major component of the architecture is the Quantum Neural Network.

In the previous study, QNNs were introduced as feature transformation mechanisms capable of generating alternative representations of financial data. The same principle applies here, but within a more sophisticated setting.

Instead of feeding raw observations directly into the DQN, selected features are first processed through parameterized quantum circuits.

The quantum layer performs a transformation of the input data into a quantum state space. Through rotations, entangling operations, and measurements, the circuit generates new representations that are subsequently incorporated into the learning process.

The rationale is straightforward.

Financial data often exhibit nonlinear interactions, hidden dependencies, and complex structures that are difficult to capture using conventional representations alone.

Quantum state spaces provide an additional representational layer through which these structures may be examined.

Whether this ultimately leads to measurable predictive advantages remains a subject for future empirical investigation. Nevertheless, the architecture demonstrates that quantum computational ideas can be incorporated naturally into broader learning systems.

More importantly, it reinforces a recurring theme throughout this paper: quantum computation is fundamentally about representation.

The QNN does not replace reinforcement learning.

It enriches the information available to the reinforcement learner.

7.5 The Agentic Layer

The most distinctive feature of the architecture is the addition of an agentic layer.

Many artificial intelligence systems operate as isolated models. Inputs are provided, outputs are generated, and the process ends there.

Real-world research rarely works in this manner.

Researchers continuously challenge assumptions, compare alternatives, and evaluate competing explanations.

The agentic framework attempts to replicate aspects of this behavior.

Multiple agents perform specialized roles within the system. Some agents generate hypotheses. Others evaluate evidence. Additional agents monitor consistency, detect anomalies, enforce governance requirements, or challenge conclusions produced elsewhere in the architecture.

The objective is not to create a single all-knowing model.

The objective is to create a structured process through which ideas are generated, tested, and refined.

This design philosophy reflects developments in scientific discovery systems, where progress increasingly emerges from coordinated interactions among specialized components rather than from isolated algorithms.

7.6 Governance and Scientific Discovery

A particularly important aspect of the study was the emphasis on governance.

Financial institutions operate within highly regulated environments. Models influence investment decisions, risk assessments, and resource allocations. Consequently, transparency and accountability are essential.

For this reason, the architecture was designed according to a governance-first philosophy.

Every stage of the discovery process can be audited.

Hypotheses are recorded.

Experimental results are documented.

Model outputs are evaluated against predefined criteria.

Alternative explanations are considered before decisions are accepted.

This approach is closely aligned with the scientific method.

A scientific investigation does not begin by assuming that a hypothesis is correct. Instead, hypotheses are treated as candidates that must survive repeated attempts at refutation.

The same principle applies here.

Signals generated by machine learning systems are treated as hypotheses rather than truths.

The architecture continuously tests, challenges, and refines its conclusions.

This governance framework is particularly important when quantum components are introduced. Quantum models can appear opaque to practitioners unfamiliar with their underlying mathematics. Embedding them within a transparent discovery process helps ensure that their contributions remain interpretable and subject to scrutiny.

7.7 A Unified Framework for Financial Intelligence

The significance of this study extends beyond the specific architecture that was implemented.

At a conceptual level, it represents the convergence of several themes explored throughout the collection.

The root-finding primer introduced the idea of state-space representations.

Portfolio optimization demonstrated how financial objectives can be encoded into alternative mathematical frameworks.

Shor’s algorithm illustrated the power of representational transformation.

Quantum neural networks showed how quantum state spaces can contribute to learning.

The agentic architecture brings these ideas together.

Optimization, learning, exploration, representation, and governance become components of a single discovery ecosystem.

The result is not merely a predictive model but a framework for financial intelligence.

This framework remains highly experimental, and many practical challenges remain unresolved. Nevertheless, it offers an intriguing glimpse into a future in which quantum computation, machine learning, and autonomous reasoning systems operate collaboratively.

Most importantly, it reinforces the central lesson of this entire collection. The true promise of quantum computational thinking may not lie exclusively in faster computation. Its deeper contribution is the ability to transform complex problems into representations where useful structures

become visible. In an agentic environment, these representations become tools not only for solving problems but also for discovering entirely new ways of understanding them.

8 Synthesis: A Unified Perspective on Quantum Computational Thinking

The title of this manuscript refers to the essential contributions of quantum computing to finance. Having completed the journey through the conceptual primer and four supported companion studies, we are now in a position to answer that question directly.

The answer is perhaps different from what many readers might initially expect.

Quantum computing is often presented as a technological revolution whose primary significance derives from the possibility of faster computation. Popular discussions frequently emphasize the prospect of solving optimization problems more efficiently, accelerating machine learning, breaking cryptographic systems, or performing calculations that would be impractical on conventional computers. While these possibilities have motivated enormous interest in quantum technologies, they do not fully capture the deeper intellectual contribution of quantum computational thinking.

The most enduring contribution of quantum computing may not be speed.

It may be representation.

Throughout the conceptual primer and four studies examined in this manuscript, a common pattern has emerged repeatedly. Problems that initially appear difficult become easier to understand once they are expressed within a different representational framework. In some cases, the new representation exposes hidden structure. In others, it reveals relationships that were previously obscured. Sometimes it transforms an apparently intractable problem into one that can be addressed through a completely different set of tools.

This observation is not unique to quantum computing. In fact, it reflects one of the most fundamental principles of scientific progress.

Human beings rarely solve complex problems by confronting them directly. Instead, they develop representations that make the problems easier to understand.

Mathematics itself is a language of representation.

Geometry represents spatial relationships.

Algebra represents patterns of transformation.

Calculus represents continuous change.

Probability theory represents uncertainty.

Statistics represents information extracted from data.

Machine learning represents structure hidden within observations.

Quantum computing should be understood as belonging to this same intellectual tradition.

Its significance lies not merely in performing calculations but in providing new ways of representing information, relationships, and possibilities.

The conceptual primer introduced this idea through the concept of the quantum state. At a superficial level, a quantum state may appear to be simply a mathematical object used to describe physical systems. From a computational perspective, however, it represents something much more important. It provides a framework through which multiple possibilities can coexist within a single representation.

This idea is profoundly relevant to finance.

Financial decision-making rarely involves certainty. Portfolio managers evaluate alternative allocations. Traders assess competing market scenarios. Risk managers analyze multiple possible futures. Researchers compare rival hypotheses regarding the behavior of economic systems.

In each case, the challenge involves reasoning about possibilities.

Quantum state representations offer a language for thinking about such possibilities in a structured manner.

The first supported study extended this perspective into optimization. Portfolio construction was transformed into an energy minimization problem. Candidate allocations became states within a landscape. Constraints became penalties. Optimization became a process of navigating structure rather than merely searching through alternatives.

The lesson extends far beyond portfolio management.

Many difficult problems become more manageable once they are represented within a landscape that reveals their underlying geometry. What appears complicated in one representation may become intuitive in another.

This observation leads naturally to the second supported study and perhaps the most powerful lesson of the entire collection.

Shor's algorithm is frequently described as a factorization algorithm. In reality, it is an algorithm about transformation.

The algorithm succeeds because it refuses to solve the original problem directly.

Factorization becomes period detection.

Period detection becomes phase estimation.

Phase estimation becomes frequency analysis.

At each stage, the representation changes.

The original problem gradually disappears and is replaced by a sequence of alternative descriptions that expose structure that was previously invisible.

This idea has implications far beyond quantum computing.

Many of the most important breakthroughs in science, mathematics, and finance arise from discovering better representations rather than better calculations. The ability to reformulate a problem is often more valuable than the ability to solve the original formulation.

For financial professionals, this may be one of the most important lessons of quantum computational thinking.

The third supported study introduced machine learning and demonstrated that representation plays an equally central role in prediction. Modern neural networks succeed because they learn use-

ful representations of data. Raw observations are transformed into progressively more informative feature spaces. Prediction improves because structure becomes easier to identify.

Quantum neural networks extend this principle into quantum state spaces.

Whether or not future research ultimately demonstrates consistent performance advantages is, in some respects, a secondary question. The more immediate lesson is conceptual. Quantum computation expands the set of representational tools available to the practitioner. It encourages exploration of alternative feature spaces, alternative transformations, and alternative descriptions of data.

The fourth supported study extended these ideas further still. Optimization, transformation, and learning were combined within an agentic discovery architecture. The resulting framework introduced a perspective that may prove increasingly important in the coming decades.

Discovery itself can be viewed as a computational process.

Researchers formulate hypotheses.

Hypotheses are tested.

Evidence is evaluated.

Alternative explanations are considered.

The process repeats.

Knowledge emerges gradually through structured exploration.

The agentic framework examined in the final study formalizes this process. Quantum representations, machine learning systems, reinforcement learning agents, and governance mechanisms become components of a broader architecture designed to support discovery rather than merely prediction.

This progression reveals a remarkable intellectual continuity across the conceptual primer and four studies.

The collection begins with representation.


Representation enables optimization.

Optimization motivates transformation.

Transformation supports learning.

Learning contributes to discovery.

The progression can therefore be summarized as follows:



At every stage, representation remains the central organizing principle.

This observation suggests that the most important contribution of quantum computing to finance may not be a particular algorithm, hardware platform, or computational advantage. Instead, it may be a new framework for thinking about complexity itself.

Financial systems are among the most complex systems created by human beings. They are adaptive, nonlinear, interconnected, and continuously evolving. Traditional analytical tools remain extraordinarily valuable, but they are often challenged by the scale and complexity of modern financial environments.

Quantum computational thinking introduces an alternative perspective.

Rather than asking how a problem can be solved within its current representation, it encourages us to ask whether the representation itself should be changed.

This question is surprisingly powerful.

Sometimes a better representation reveals a hidden pattern.

Sometimes it reveals a simpler optimization landscape.

Sometimes it exposes a useful transformation.

Sometimes it enables a new learning architecture.

Sometimes it changes the entire nature of the problem.

This idea constitutes the unifying theme of the present manuscript.

The four supported companion studies and the conceptual primer were intentionally chosen to illustrate different manifestations of the same underlying principle. Root finding, portfolio optimization, factorization, machine learning, and agentic discovery may appear unrelated on the surface. Yet each example demonstrates the value of transforming a problem into a representation in which useful structure becomes visible.

For this reason, the ultimate lesson of quantum computing may be neither physical nor technological.

It may be epistemological.

Quantum computational thinking teaches us that understanding often emerges from representation. The way in which a problem is described can be as important as the method used to solve it.

This insight has guided scientific progress for centuries. Quantum computing extends that tradition into a new computational domain.

For financial professionals, the practical implication is straightforward. One does not need to become a quantum physicist to benefit from quantum computational thinking. The essential ideas can already inform how we approach optimization, learning, forecasting, risk management, and decision-making. They encourage us to search for better representations, more revealing transformations, and richer descriptions of complex systems.

In that sense, quantum computing should be viewed not merely as a new technology but as a new intellectual framework. The true promise of that framework lies in its ability to help us think differently about problems whose complexity might otherwise appear overwhelming.

That, ultimately, is the essential contribution of quantum computing to finance.

9 Companion Computational Laboratories

The educational philosophy underlying this collection combines conceptual discussion with practical experimentation. Each companion study is supported by one or more executable Colab notebooks that allow readers to explore the computational ideas presented throughout the manuscript.

The notebooks are not intended to serve as production systems. Rather, they function as educational laboratories through which readers can interact directly with quantum-inspired computational frameworks, inspect intermediate results, and develop intuition regarding the concepts discussed in the text.

Readers are encouraged to study the manuscripts and notebooks together.

9.1 Conceptual Primer: A Simple Quantum Approach to Finding the Root of a Function

This introductory root-finding material is included as a conceptual primer rather than as a supported companion paper. No dedicated paper or notebook accompanies it in the collection.

The concepts introduced here are intentionally simple and serve primarily as a conceptual foundation for the four supported companion studies. Readers should focus on developing intuition regarding quantum states, superposition, interference, and measurement before proceeding to the subsequent computational laboratories.

9.2 Companion Study 1: Quantum Portfolio Optimization

Companion Notebook Repository:

https://drive.google.com/drive/folders/1YXWGHZ9e0_oSz6Dfx4bq-a5-jHLZzYmY?usp=sharing

These notebooks demonstrate how portfolio allocation problems can be represented through Hamiltonian formulations and explored using quantum-inspired optimization frameworks.

The notebooks complement the portfolio optimization manuscript by illustrating how candidate allocations become computational states within an optimization landscape and how financial constraints can be encoded within alternative representations.

9.3 Companion Study 2: A Governance-First Agentic Method for Teaching and Implementing Shor’s Algorithm: From Modular Period Finding to Auditable Quantum Workflow

Companion Notebook:

<https://drive.google.com/file/d/1nSatQBozD6dMoiQhkIkRWN486Q0HHvqh/view?usp=sharing>

This notebook provides an educational implementation of Shor’s algorithm using an agentic architecture.

The workflow decomposes the algorithm into specialized components responsible for state preparation, modular arithmetic, phase estimation, quantum Fourier transformation, and result interpretation.

The notebook complements the manuscript by transforming one of the most celebrated quantum algorithms into a transparent and interactive educational experience.

9.4 Companion Study 3: Hybrid Quantum Neural Networks for Financial Prediction

Companion Notebook Repository:

<https://drive.google.com/drive/folders/13QTfgWDSUPpBtzlF93C2QbiXf0yUwk67?usp=sharing>

These notebooks explore hybrid machine learning architectures combining classical convolutional neural networks with parameterized quantum circuits.

The computational laboratories allow readers to examine how quantum-inspired representations can be incorporated into predictive systems and how alternative feature spaces can be constructed through hybrid quantum-classical models.

9.5 Companion Study 4: An Agentic DQN-DNN-QNN Architecture for Financial Discovery

Companion Notebook:

<https://drive.google.com/file/d/1s5r4C4EHGN7U1BlzA473C-7DrGu90mTY/view?usp=sharing>

This notebook implements the most sophisticated architecture within the collection.

Deep Q-Networks, Deep Neural Networks, Quantum Neural Networks, reinforcement learning mechanisms, and governance-first agentic workflows are integrated into a unified framework designed to support financial discovery and autonomous research.

The notebook serves as the computational culmination of the educational journey developed throughout the manuscript.

9.6 Additional Computational Laboratory: Agentic Quantum Optimization

Companion Notebook:

<https://drive.google.com/file/d/1oenIHLqH-OVFh6NPYqIn8DQv1PKcgS3A/view?usp=sharing>

This notebook extends the optimization themes developed in the portfolio optimization study through a governance-first agentic architecture.

The laboratory illustrates how optimization workflows can be decomposed into specialized computational agents responsible for representation, evaluation, optimization, validation, and interpretation.

As such, it serves as a bridge between the optimization concepts introduced in the portfolio study and the broader discovery architecture presented in the final manuscript.

9.7 Recommended Study Sequence

Readers are encouraged to follow the conceptual primer first, followed by the four supported companion studies:

1. Conceptual primer: A Simple Quantum Approach to Finding the Root of a Function
2. Quantum Portfolio Optimization
3. Agentic Quantum Optimization Notebook
4. A Governance-First Agentic Method for Teaching and Implementing Shor’s Algorithm: From Modular Period Finding to Auditable Quantum Workflow
5. Hybrid Quantum Neural Networks for Financial Prediction
6. An Agentic DQN-DNN-QNN Architecture for Financial Discovery

This sequence mirrors the intellectual progression developed throughout the collection and provides a gradual transition from elementary quantum concepts to advanced agentic computational architectures.

Together, the manuscripts and companion notebooks form an integrated educational ecosystem designed to help financial professionals develop intuition regarding quantum computational thinking and its potential applications within modern finance.

10 Limitations and Future Directions

Every educational work must make choices regarding scope, depth, and audience. The present manuscript is no exception. The four supported companion studies summarized throughout this

paper were designed with a specific objective in mind: to provide financial professionals with an accessible introduction to the most important conceptual contributions of quantum computing and quantum-inspired computation.

As a consequence, the collection deliberately prioritizes intuition over mathematical completeness, conceptual understanding over formal proof, and educational clarity over algorithmic sophistication. These choices have enabled the development of a coherent learning pathway, but they also define the boundaries of the work. Understanding those boundaries is important both for appreciating what the collection accomplishes and for identifying opportunities for future exploration.

10.1 An Educational Rather Than Experimental Contribution

The first limitation of the collection is also one of its defining characteristics.

The manuscript is educational rather than experimental.

The primary objective is not to establish new theoretical results, introduce novel quantum algorithms, or demonstrate empirical superiority over existing classical methods. Instead, the collection seeks to help readers understand the principles underlying quantum computational thinking and to illustrate how those principles may be relevant to finance.

Throughout the conceptual primer and four studies, quantum concepts are introduced through carefully selected examples designed to emphasize intuition and accessibility. The examples are intentionally chosen to illuminate ideas rather than maximize technical sophistication.

Consequently, the manuscript should be viewed as a conceptual bridge between finance and quantum computing rather than as a contribution to the frontier of quantum algorithm research.

The success of the collection should therefore be evaluated according to its educational effectiveness rather than its ability to establish new scientific results.

10.2 Simplified Computational Environments

A second limitation concerns the nature of the examples themselves.

Real-world financial systems are extraordinarily complex. Markets evolve continuously. Participants adapt their behavior. Data are noisy, incomplete, and often non-stationary. Practical investment decisions are influenced by transaction costs, liquidity constraints, regulation, taxation, market impact, and operational considerations.

No educational example can fully capture this complexity.

For this reason, the studies presented throughout the collection rely on simplified environments designed to isolate specific concepts. Portfolio optimization examples focus on illustrating optimization landscapes. The root-finding primer emphasizes state-space thinking and is explicitly not counted as a supported companion paper. Shor’s algorithm is presented through small-scale examples suitable for educational purposes. Machine learning architectures employ synthetic or controlled datasets intended to facilitate interpretation.

These simplifications are intentional.

The objective is to make the computational ideas visible. Introducing the full complexity of institutional financial environments at the outset would obscure the conceptual lessons that the studies are intended to communicate.

Readers should therefore view the examples as educational laboratories rather than complete financial implementations.

10.3 The Current State of Quantum Hardware

The practical deployment of many quantum algorithms remains constrained by the present state of quantum computing technology.

Contemporary quantum hardware continues to operate within what is commonly described as the Noisy Intermediate-Scale Quantum (NISQ) era. Existing systems are characterized by limited qubit counts, finite coherence times, noise, gate errors, and restrictions on circuit depth. While progress in hardware development has been substantial, large-scale fault-tolerant quantum computers capable of supporting industrial financial workloads have not yet become widely available.

This reality influences the nature of the examples presented in the collection.

Many of the notebooks and educational implementations emphasize quantum-inspired computational thinking rather than direct deployment on advanced quantum hardware. In doing so, they focus attention on representational frameworks, optimization concepts, learning architectures, and discovery processes that remain valuable regardless of the pace of hardware development.

Indeed, one of the central themes of this manuscript is that the educational value of quantum computing extends beyond the hardware itself.

10.4 No Claim of Quantum Advantage

A particularly important limitation concerns the interpretation of performance.

The examples presented throughout this collection should not be interpreted as demonstrations of quantum advantage.

Establishing quantum advantage requires rigorous benchmarking against state-of-the-art classical methods under carefully controlled experimental conditions. Such investigations involve extensive empirical analysis and remain active areas of research across academia and industry.

The present collection does not attempt to resolve these questions.

Its purpose is different.

The studies are designed to illustrate concepts, computational architectures, and representational frameworks. They seek to explain how quantum computational thinking operates rather than to establish definitive claims regarding practical superiority.

This distinction is essential because it helps place the work within its proper context. The value of the collection lies in understanding new ways of representing and reasoning about problems rather than in proving that one computational paradigm is universally superior to another.

10.5 Interpretability and Governance

As the collection progresses, the computational architectures become increasingly sophisticated. Hybrid neural networks, reinforcement learning systems, and agentic discovery frameworks introduce significant representational flexibility, but they also introduce challenges related to interpretability.

Complex systems can become difficult to understand even when implemented entirely within classical computational frameworks. The addition of quantum representations may further complicate interpretation, particularly for practitioners who lack formal training in quantum information theory.

This challenge highlights the importance of governance.

One of the recurring themes of the collection is that increasingly sophisticated computational systems require increasingly sophisticated mechanisms for oversight, validation, and auditability. The governance-first philosophy adopted in the later studies reflects the belief that transparency and accountability must remain central considerations as computational architectures continue to evolve.

Future research will need to continue exploring how powerful learning systems can remain interpretable, controllable, and trustworthy.

10.6 Future Directions

Although the collection is primarily educational, it naturally suggests a variety of promising directions for future investigation.

One avenue involves extending the optimization frameworks to larger and more realistic financial environments. Portfolio optimization examples could incorporate broader asset universes, dynamic constraints, transaction costs, and evolving market conditions.

A second direction concerns machine learning. Future work may explore richer hybrid architectures involving quantum neural networks, transformers, foundation models, and other emerging approaches to representation learning.

A third direction involves discovery systems. The agentic architectures introduced in the final study provide only an initial glimpse into the possibility of autonomous financial research environments. Future systems may incorporate larger numbers of specialized agents, more sophisticated governance frameworks, and increasingly advanced mechanisms for hypothesis generation and evaluation.

Perhaps most importantly, future work can continue expanding the educational dimension of quantum finance. As quantum computing matures, there remains a substantial need for materials capable of bridging the gap between theoretical quantum science and practical financial applications.

10.7 A Final Reflection on Scope

The limitations discussed above should not be interpreted as shortcomings of the collection. Rather, they reflect deliberate choices regarding purpose and audience.

The objective of the manuscript is not to transform readers into quantum physicists.

Nor is it to persuade them that quantum computing will immediately revolutionize every aspect of finance.

Instead, the goal is more modest and, perhaps, more durable.

The collection seeks to provide financial professionals with a conceptual framework through which they can understand the emerging language of quantum computation. It introduces a set of ideas, representations, and ways of thinking that may prove increasingly valuable as both quantum computing and artificial intelligence continue to evolve.

If readers complete the journey with a deeper appreciation for the importance of representation, transformation, learning, optimization, and discovery, then the manuscript will have achieved its primary objective.

In that sense, the work should be viewed not as a destination but as an invitation—a starting point for further exploration at the intersection of quantum computing, artificial intelligence, and modern finance.

11 Conclusion

This manuscript began with a simple question: what are the essential contributions of quantum computing to finance?

The answer that emerged throughout the four supported companion studies is both broader and more subtle than many readers might initially expect.

Quantum computing is frequently discussed in terms of computational speed, hardware development, cryptographic disruption, or the possibility of solving certain classes of problems more efficiently than classical computers. These topics are important and will undoubtedly continue to shape both academic research and industrial investment in the years ahead. Yet the educational journey presented throughout this collection suggests that the significance of quantum computing extends beyond any particular device, algorithm, or technological milestone.

At its core, quantum computing introduces a new way of thinking about problems.

The central theme connecting every chapter of this manuscript has been representation. Again and again, we observed that problems become easier to understand when they are expressed in representations that reveal their underlying structure. In some cases, these representations support optimization. In others, they expose hidden patterns, enable useful transformations, facilitate learning, or create entirely new frameworks for discovery. Regardless of the specific application, the same lesson repeatedly emerged: the way a problem is represented can be as important as the method used to solve it.

This observation provides a useful perspective for financial professionals.

Finance is fundamentally concerned with complexity. Markets are adaptive systems populated by heterogeneous participants. Economic relationships evolve through time. Information arrives continuously. Decisions must be made under uncertainty. Portfolio managers, traders, researchers, regulators, and risk managers all confront environments in which the number of possible states, scenarios, and interactions can quickly become overwhelming.

Traditional quantitative methods have provided remarkably effective tools for navigating this complexity. Portfolio theory, stochastic processes, econometrics, optimization, machine learning, and artificial intelligence have transformed the practice of finance over the past century. Quantum computational thinking should not be viewed as a replacement for these methods. Rather, it should be understood as an extension of the broader scientific tradition of developing increasingly powerful representational frameworks.

The conceptual primer and the four supported companion studies illustrate this progression from multiple perspectives.

The conceptual primer introduced the notion of the quantum state and demonstrated that computation can be viewed as the evolution of possibilities within a structured state space. The objective was not to solve difficult numerical problems but to establish a new intuition regarding how information may be represented; it is not counted as a separate supporting paper.

The first supported study showed how portfolio optimization can be interpreted through the lens of energy landscapes. Allocation decisions became candidate states and constraints became structural features of the optimization environment. The emphasis was not on replacing classical portfolio theory but on illustrating how alternative representations can provide new perspectives on familiar problems.


The second supported study explored Shor’s algorithm and highlighted one of the most profound lessons in computational science: difficult problems often become manageable when transformed into different problems. Factorization became period detection. Period detection became phase estimation. The algorithm demonstrated the extraordinary power of representational transformation.

The third supported study shifted attention toward learning. Hybrid quantum neural networks illustrated how quantum state spaces may serve as environments for representation learning and feature construction. Whether future research ultimately establishes widespread performance advantages is, in some respects, secondary to the conceptual lesson. Quantum computing expands the collection of representational tools available to researchers and practitioners.

The fourth supported study completed the journey by introducing agentic discovery architectures. Optimization, transformation, learning, governance, and evaluation were integrated into a structured process of computational discovery. The resulting framework suggested how future financial research environments may evolve as artificial intelligence and quantum-inspired methods continue to mature.

Viewed individually, these studies address different topics.

Viewed collectively, they reveal a coherent intellectual progression:



This progression captures the conceptual architecture of the entire collection.

The journey begins with representation.

Representation enables optimization.

Optimization motivates transformation.

Transformation supports learning.

Learning contributes to discovery.

At every stage, representation remains the central organizing principle.

For this reason, the most important contribution of quantum computing to finance may not be a particular algorithm or hardware platform. It may instead be the introduction of a new language for reasoning about complexity.

This language encourages practitioners to ask a different set of questions.

Rather than asking only how a problem can be solved, one may ask whether the problem has been represented appropriately.

Rather than asking only how a model can be improved, one may ask whether a more revealing representation exists.

Rather than asking only how computation can be accelerated, one may ask whether transformation can expose structure that was previously hidden.

These questions are valuable regardless of whether a computation ultimately takes place on a classical computer, a quantum computer, or some future hybrid architecture.

Indeed, one of the central conclusions of this manuscript is that many of the educational benefits of quantum computing are available today. They do not depend upon the arrival of large-scale fault-tolerant quantum hardware. They arise from the conceptual framework itself. The ideas of state-space representation, transformation, interference, optimization, learning, and discovery can already influence how financial professionals think about their work.

This observation also explains the importance of the companion computational notebooks. Throughout the collection, theoretical discussion has been paired with executable computational laboratories. The notebooks allow readers to move beyond abstract descriptions and observe the behavior of quantum-inspired architectures directly. They transform concepts into experiences and provide a practical pathway through which intuition can be developed.

The educational philosophy underlying the collection is therefore simple. Quantum computing should not be approached exclusively as a branch of physics, nor solely as a future technology. It should also be viewed as a framework for thinking. By engaging with the notebooks, experimenting with the models, and exploring the companion manuscripts, readers are invited to participate actively in that process of intellectual discovery.

It is equally important to recognize what this manuscript has not attempted to do. The collection does not claim to establish quantum advantage. It does not attempt to prove that quantum methods will replace classical approaches. Nor does it seek to predict the precise trajectory of future quantum hardware. Such questions remain active areas of research and will undoubtedly continue to evolve over the coming decades.

The purpose of this work has been more modest and, perhaps, more enduring.

Its objective has been to create a bridge.

On one side of that bridge lies the rapidly evolving world of quantum computing. On the other lies the practical world of financial decision-making. Between them stands a community of professionals seeking to understand how emerging computational paradigms may influence the future of their discipline.

The conceptual primer and four studies presented throughout this manuscript represent one possible pathway across that bridge.

They demonstrate that quantum computing is not merely a collection of algorithms. It is a framework for representing possibilities, transforming problems, learning from data, and organizing discovery. It provides a new perspective on complexity and a new vocabulary through which complex systems may be understood.

Ultimately, the future significance of quantum computing in finance will not be determined solely by advances in hardware or by the success of individual algorithms. It will also depend upon whether practitioners learn to think differently about the problems they face.

If this collection encourages readers to view optimization, forecasting, learning, and decision-making through a richer representational lens, then it will have achieved its primary objective.

In that sense, the journey presented throughout these pages should not be viewed as a conclusion.

It should be viewed as a beginning.

A Concept Map of Quantum Computational Thinking

This appendix summarizes the conceptual architecture developed throughout the manuscript. Rather than presenting quantum computing as a collection of isolated algorithms or mathematical techniques, the collection has emphasized a unified perspective in which quantum computational thinking is understood as a framework for representing, transforming, learning from, and discovering structure within complex systems.

The purpose of this concept map is to provide readers with a high-level overview of the in-


tellectual progression developed across the conceptual primer and the four supported companion studies.

A.1 The Central Question

The entire collection can be understood as an attempt to answer a single question:

What are the essential contributions of quantum computing to finance?

The answer developed throughout the manuscript is summarized below.



These five ideas constitute the conceptual backbone of the collection.


A.2 Level 1: Representation

The first contribution of quantum computational thinking is the introduction of alternative representational frameworks.

Traditional computation represents information through deterministic binary states.

Quantum computation represents information through state vectors composed of amplitudes, phases, and basis states.

From an educational perspective, the significance of the quantum state is not primarily physical. Its importance lies in the fact that it provides a new language for describing possibilities.



The root-finding primer was designed to establish intuition regarding this foundational concept, without being counted as a supporting companion paper.

A.3 Level 2: Optimization

Once possibilities have been represented, the next challenge is identifying desirable configurations.


The portfolio optimization study demonstrated how financial decision problems can be transformed into optimization landscapes.

Portfolio allocations become candidate states.

Constraints become penalties.

Objectives become energy functions.

Optimization becomes a search for low-energy configurations.




The central lesson is that optimization often depends as much on representation as on the search procedure itself.

A.4 Level 3: Transformation

The third level introduces one of the most important ideas in computational science.

Many difficult problems become easier when they are transformed into different problems.

Shor's algorithm provides a canonical example.



The algorithm demonstrates that the path to a solution frequently involves changing the representation of the problem rather than attacking the original formulation directly.

This principle extends far beyond factorization and serves as one of the central themes of the collection.


A.5 Level 4: Learning

Machine learning introduces a new objective.

Rather than searching directly for solutions, learning systems attempt to discover useful representations of data.

Neural networks operate by constructing progressively richer feature spaces.

Quantum neural networks extend this idea by introducing quantum state spaces as environments for representation learning.



The third supported study demonstrates how quantum computational ideas can contribute to modern predictive systems through alternative representational frameworks.

A.6 Level 5: Discovery

The final level extends optimization, transformation, and learning into a broader process of discovery.

Agentic systems do not merely generate predictions.


They formulate hypotheses.

They evaluate evidence.

They compare alternatives.

They refine explanations.


They learn from outcomes.



The fourth supported study illustrates how these components can be integrated into governance-first research architectures.

A.7 The Unified Framework

The conceptual primer and four studies collectively reveal a common computational pattern.



Each stage builds upon the previous stage.

Representation creates structure.

Transformation reveals hidden relationships.

Optimization identifies desirable configurations.


Learning extracts patterns.

Discovery generates knowledge.

The progression therefore represents a general framework for reasoning about complexity.

A.8 The Relationship Between the Conceptual Primer and Four Companion Studies

The educational structure of the collection can be summarized as follows. The first item is a conceptual primer without a separate supporting paper; the remaining four items are supported companion studies:



A Governance-First Agentic Method for Teaching and Implementing Shor’s Algorithm: From
Modular Period Finding to Auditable Quantum Workflow

↓

Core Idea: Transformation

Study 3

Hybrid Quantum Neural Networks

↓

Core Idea: Learning

Study 4

Agentic DQN-DNN-QNN Architectures

↓

Core Idea: Discovery

The sequence is intentional.

Each study introduces a new layer of abstraction while preserving the conceptual foundations established by the preceding studies.

A.9 The Essential Contributions of Quantum Computing to Finance

The collection ultimately identifies five essential contributions of quantum computational thinking.

1. The ability to represent possibilities through quantum states.
2. The ability to reformulate optimization problems within alternative computational landscapes.
3. The ability to transform difficult problems into representations that reveal hidden structure.
4. The ability to construct alternative representation spaces for learning and prediction.
5. The ability to integrate optimization, learning, and transformation within broader discovery architectures.

These contributions do not depend exclusively on future hardware developments.

They arise from the conceptual framework itself.

For this reason, the most enduring contribution of quantum computing to finance may not be computational speed alone.

It may be the introduction of a richer language for understanding complexity.

This concept map summarizes that language and provides a unifying framework through which the individual studies presented throughout the manuscript can be understood as components of a single educational journey.

B Glossary of Quantum Computing Terms for Financial Professionals

The purpose of this glossary is not to provide formal mathematical definitions. Instead, it offers intuitive descriptions of the concepts introduced throughout the manuscript. The emphasis is on understanding how these ideas contribute to quantum computational thinking and why they may be relevant to finance.

Agentic Architecture	A computational framework composed of multiple specialized agents that collaborate to perform tasks such as optimization, learning, evaluation, governance, and discovery. In the context of this manuscript, agentic architectures represent the highest level of abstraction in the educational progression.
Amplitude	A numerical quantity associated with a quantum state. Amplitudes determine the probability of observing particular outcomes after measurement. Conceptually, amplitudes indicate the relative importance of different possibilities within a quantum representation.
Basis State	One of the fundamental states used to describe a quantum system. In computational terms, basis states represent the individual possibilities from which more complex quantum states are constructed.
Classical Computation	The traditional computational framework based on deterministic bits that take values of either 0 or 1. Quantum computation extends rather than replaces this framework.
Controlled Operation	A quantum operation whose behavior depends on the state of another qubit. Controlled operations allow information contained in one part of a quantum system to influence another part of the system.
Deep Neural Network (DNN)	A machine learning architecture composed of multiple layers of artificial neurons. DNNs learn useful representations of data and serve as one of the foundational components of modern artificial intelligence.
Deep Q-Network (DQN)	A reinforcement learning architecture that combines neural networks with decision-making algorithms. DQNs learn policies through interaction with an environment and were incorporated into the discovery framework discussed in the final study.
Discovery	The process of generating, evaluating, refining, and validating hypotheses. Within the educational framework developed throughout this manuscript, discovery represents the highest level of computational capability.
Energy Landscape	A representation of an optimization problem in which candidate solutions are associated with energy values. Desirable solutions correspond to low-energy configurations.

Entanglement	A relationship between quantum systems in which the state of one system cannot be described independently of the other. Although often viewed as a physical phenomenon, entanglement can also be understood as a powerful mechanism for representing relationships between variables.
Feature Representation	A transformed description of data that makes useful patterns easier to identify. Representation learning is one of the central ideas connecting machine learning and quantum computation.
Governance-First Architecture	A design philosophy emphasizing transparency, auditability, validation, and human oversight. Governance-first principles become increasingly important as computational systems grow more complex and autonomous.
Hamiltonian	A mathematical object used to represent the structure of an optimization problem. In the portfolio optimization study, the Hamiltonian serves as the mechanism through which objectives and constraints are encoded.
Hybrid Quantum-Classical Model	A computational architecture that combines classical machine learning methods with quantum-inspired or quantum computational components.
Interference	The process through which quantum amplitudes combine. Interference allows certain possibilities to become more prominent while others become less prominent, making it a key mechanism in many quantum algorithms.
Learning	The process of constructing useful representations from data. In the context of this manuscript, learning represents the fourth stage in the progression from representation to discovery.
Measurement	The process through which a quantum state produces an observable outcome. Measurement converts a collection of possibilities into a specific result.
Optimization	The process of identifying desirable solutions among many alternatives. Optimization represents one of the most important applications of quantum-inspired computational thinking in finance.
Parameterized Quantum Circuit (PQC)	A quantum circuit containing adjustable parameters that can be trained in a manner analogous to neural network weights.
Phase	A property of a quantum state that influences how amplitudes combine through interference. Although not directly observable, phase often contains important computational information.

Phase Estimation	A quantum computational procedure used to extract phase information from a quantum system. Phase estimation plays a central role in Shor’s algorithm.
Quantum Computational Thinking	The central theme of this manuscript. It refers to the use of quantum-inspired concepts such as state representations, transformation, interference, optimization, learning, and discovery as tools for reasoning about complex problems.
Quantum Discovery	The integration of representation, optimization, transformation, learning, and governance into a unified framework capable of supporting autonomous research and decision-making.
Quantum Fourier Transform (QFT)	A quantum transformation that reveals frequency and periodicity information contained within a quantum state. The QFT plays a critical role in Shor’s algorithm and serves as an example of representational transformation.
Quantum Neural Network (QNN)	A machine learning architecture that incorporates quantum state representations and parameterized quantum circuits into the learning process.
Quantum Optimization	The use of quantum-inspired representations and computational frameworks to solve optimization problems.
Quantum State	A mathematical representation of possibilities. Throughout this manuscript, the quantum state serves as the foundational concept from which all other ideas emerge.
Quantum Transformation	The process of reformulating a problem into a representation that exposes useful structure. Shor’s algorithm provides the most important example discussed in the collection.
Qubit	The fundamental unit of quantum information. While a classical bit represents either 0 or 1, a qubit can represent combinations of possibilities through superposition.
Reinforcement Learning	A machine learning framework in which an agent learns through interaction with an environment and feedback from outcomes.
Representation	The central concept of the manuscript. A representation is a way of describing a problem. Much of the power of quantum computational thinking arises from discovering representations that reveal useful structure.
Shor’s Algorithm	A quantum algorithm famous for integer factorization. Educationally, its most important lesson is that difficult problems can become tractable when transformed into more revealing representations.

State Space	The collection of all possible configurations of a system. Quantum computational thinking frequently begins by constructing and exploring state spaces.
Superposition	The ability of a quantum state to represent multiple possibilities simultaneously. Superposition provides the foundation for quantum state-space representations.
Transformation	The process of converting a problem into an alternative representation that reveals useful structure. Transformation constitutes one of the central themes of the collection.
Unitary Transformation	A reversible operation that changes a quantum state while preserving its overall structure. Quantum algorithms are built through sequences of such transformations.

A Final Observation

A recurring theme throughout this glossary is that many quantum concepts can be interpreted as concepts about representation. Quantum states represent possibilities. Transformations represent changes of perspective. Optimization represents the search for desirable configurations. Learning represents the construction of useful descriptions of data. Discovery represents the integration of these capabilities into a coherent process of knowledge generation.

For this reason, readers may find it useful to remember a single guiding principle:

The most important contribution of quantum computational thinking is not computation alone. It is the ability to represent complexity in new and useful ways.

This principle provides the conceptual thread connecting the conceptual primer, all four supported studies, and the broader educational journey developed throughout the manuscript.

C Educational Roadmap for Readers

This manuscript was designed as an educational journey rather than a conventional collection of independent studies. The conceptual primer introduces the state-space intuition, and the four companion papers then build progressively upon one another, introducing increasingly sophisticated ideas while maintaining a consistent conceptual thread.

For this reason, readers are encouraged to approach the material sequentially.

The objective of the roadmap presented in this appendix is to help readers understand what they should learn at each stage of the journey and how the individual studies contribute to a broader understanding of quantum computational thinking.

C.1 The Educational Philosophy

The collection was designed around a simple pedagogical principle.

Most introductions to quantum computing begin with physics.

This collection begins with problems.

Rather than asking readers to master the mathematics of quantum mechanics before exploring applications, the studies introduce quantum concepts through computational and financial examples. The emphasis is on intuition, representation, and problem-solving rather than formal derivation.

The educational progression therefore follows the logic of computational thinking rather than the historical development of quantum physics.

The ultimate goal is not to transform financial professionals into quantum physicists. The goal is to help them understand how quantum-inspired ideas can influence the way complex problems are represented, transformed, optimized, learned, and ultimately solved.

C.2 Stage One: Learning to Think in State Spaces

Conceptual Primer:

A Simple Quantum Approach to Finding the Root of a Function

Note: This stage has no separate supporting paper in the companion collection.

The first stage introduces the foundational concept of the entire collection: the quantum state.

Readers are encouraged to focus less on the mechanics of the algorithm and more on the underlying shift in perspective. The key lesson is that computation can be viewed as the evolution of a system through a structured space of possibilities.

At this stage, readers should become comfortable with the following ideas:

- Quantum states
- Basis states
- Superposition
- Amplitudes
- Measurement
- State-space representations

The educational objective is not mathematical mastery.

The objective is to begin thinking about problems as collections of possibilities rather than isolated outcomes.

C.3 Stage Two: Learning to Think in Optimization Landscapes

Companion Study:

Quantum Portfolio Optimization

The second stage extends state-space thinking into optimization.

Readers learn how financial decision problems can be represented as landscapes in which candidate solutions occupy different positions and possess different energy levels.

The most important lesson is that optimization depends heavily upon representation.

At this stage, readers should focus on understanding:

- Hamiltonian representations
- Energy landscapes
- Constraints as penalties
- Candidate solution states
- Optimization as search

The objective is to recognize that alternative representations can make difficult optimization problems easier to analyze.

C.4 Stage Three: Learning the Power of Transformation

Companion Study:

A Governance-First Agentic Method for Teaching and Implementing Shor's Algorithm: From Modular Period Finding to Auditable Quantum Workflow

The third stage introduces what is arguably the most profound lesson of the entire collection.

Many difficult problems become easier when transformed into different problems.

Readers should approach Shor's algorithm not primarily as a factorization algorithm but as an example of representational transformation.

The essential concepts introduced at this stage include:

- Periodicity
- Phase estimation
- Controlled operations
- Quantum Fourier Transform
- Problem reformulation

The key educational takeaway is simple:

Sometimes the most effective way to solve a problem is to stop trying to solve the original problem.

Instead, one seeks a representation in which hidden structure becomes visible.

C.5 Stage Four: Learning Through Representation

Companion Study:

Hybrid Quantum Neural Networks for Financial Prediction

The fourth stage introduces machine learning.

By this point, readers have already encountered state representations, optimization, and transformation. The next step is understanding how computational systems learn useful representations from data.

Readers should focus on:

- Feature construction
- Representation learning
- Neural networks
- Parameterized quantum circuits
- Hybrid quantum-classical architectures

The objective is not to determine whether quantum neural networks outperform classical neural networks.

Instead, the goal is to understand how quantum state spaces can serve as alternative environments for representation learning.

C.6 Stage Five: Learning Through Discovery

Companion Study:

An Agentic DQN-DNN-QNN Architecture for Financial Discovery

The final stage integrates all previous concepts.

Representation.

Optimization.

Transformation.

Learning.

These capabilities are combined within a governance-first discovery architecture.

Readers should focus on understanding:

- Reinforcement learning
- Deep Q-Networks
- Agentic systems
- Governance mechanisms

- Autonomous discovery processes


The educational objective is to recognize that discovery itself can be viewed as a computational activity.

The system no longer seeks a single answer.

Instead, it continuously generates, evaluates, and refines hypotheses.

C.7 The Complete Journey

The intellectual progression of the collection can be summarized as follows:



Each stage depends upon the previous stage.

Representation creates possibilities.

Optimization identifies desirable possibilities.

Transformation reveals hidden structure.

Learning extracts patterns.

Discovery generates knowledge.

Together, these stages form the conceptual framework that this manuscript refers to as quantum computational thinking.

C.8 Recommended Reading Sequence

Readers are encouraged to follow the conceptual primer and supported studies in the following order:

1. Conceptual primer: A Simple Quantum Approach to Finding the Root of a Function
2. Quantum Portfolio Optimization
3. A Governance-First Agentic Method for Teaching and Implementing Shor's Algorithm: From Modular Period Finding to Auditable Quantum Workflow
4. Hybrid Quantum Neural Networks for Financial Prediction

5. An Agentic DQN-DNN-QNN Architecture for Financial Discovery

This sequence mirrors the conceptual progression developed throughout the manuscript and minimizes the need for prior knowledge.

C.9 Recommended Notebook Sequence

The companion computational laboratories should be explored in parallel with the manuscripts.

A suggested sequence is:

1. Quantum Portfolio Optimization
2. Agentic Quantum Optimization
3. A Governance-First Agentic Method for Teaching and Implementing Shor's Algorithm: From Modular Period Finding to Auditable Quantum Workflow
4. Hybrid CNN-QNN Neural Networks
5. DQN-DNN-QNN Reinforcement Learning Architecture

Readers are encouraged to modify parameters, inspect intermediate outputs, and experiment with alternative assumptions. The notebooks are intended to function as educational laboratories rather than fixed demonstrations.

C.10 The Final Learning Objective

Upon completing the conceptual primer, the four supported studies, and their companion notebooks, readers should not expect to have mastered quantum mechanics.

Nor should they expect to have become specialists in quantum algorithm design.

Instead, they should possess something that may ultimately be more valuable.

They should understand a new way of thinking about complex problems.

The most important lesson of the collection is that representation matters.

Quantum computational thinking provides a framework for representing possibilities, transforming problems, learning from data, and organizing discovery. Whether future financial systems ultimately run on classical computers, quantum computers, or hybrid architectures, these ideas will remain relevant because they address a more fundamental question:

How can complex systems be represented in ways that make useful structure visible?

That question lies at the heart of both quantum computing and finance.

It is also the question that has guided the educational journey developed throughout this manuscript.

References

- [1] Reynoso, A. *Quantum Portfolio Optimization: A Canonical Educational Framework*. Manuscript, Judge Business School, University of Cambridge, June 2026.
- [2] Reynoso, A. *A Governance-First Agentic Method for Teaching and Implementing Shor’s Algorithm: From Modular Period Finding to Auditable Quantum Workflow*. Manuscript, Judge Business School, University of Cambridge, June 2026.
- [3] Reynoso, A. *Hybrid Quantum Neural Networks for Financial Prediction*. Manuscript, Judge Business School, University of Cambridge, June 2026.
- [4] Reynoso, A. *An Agentic DQN-DNN-QNN Architecture for Financial Discovery*. Manuscript, Judge Business School, University of Cambridge, June 2026.
- [5] Feynman, R. P. *Simulating Physics with Computers*. International Journal of Theoretical Physics, 21(6–7), 467–488, 1982.
- [6] Nielsen, M. A., and Chuang, I. L. *Quantum Computation and Quantum Information*. Cambridge University Press, 2010.
- [7] Grover, L. K. *A Fast Quantum Mechanical Algorithm for Database Search*. Proceedings of the Twenty-Eighth Annual ACM Symposium on Theory of Computing, 1996.
- [8] Shor, P. W. *Polynomial-Time Algorithms for Prime Factorization and Discrete Logarithms on a Quantum Computer*. SIAM Journal on Computing, 26(5), 1484–1509, 1997.
- [9] Farhi, E., Goldstone, J., and Gutmann, S. *A Quantum Approximate Optimization Algorithm*. arXiv:1411.4028, 2014.
- [10] Markowitz, H. *Portfolio Selection*. Journal of Finance, 7(1), 77–91, 1952.
- [11] Schuld, M., Sinayskiy, I., and Petruccione, F. *The Quest for a Quantum Neural Network*. Quantum Information Processing, 13(11), 2567–2586, 2014.
- [12] Biamonte, J., Wittek, P., Pancotti, N., Rebentrost, P., Wiebe, N., and Lloyd, S. *Quantum Machine Learning*. Nature, 549, 195–202, 2017.
- [13] Sutton, R. S., and Barto, A. G. *Reinforcement Learning: An Introduction*. Second Edition, MIT Press, 2018.
- [14] Russell, S. *Human Compatible: Artificial Intelligence and the Problem of Control*. Penguin Books, 2021.